

SLR210

Project: Obstruction-Free Consensus and Paxos

The goal of this project is to get an initial experience in designing a fault-tolerant distributed system. Here we focus on a state-machine replicated system build atop a consensus abstraction.

1 Specification

An obstruction-free consensus (OFC) algorithm exports one operation *propose*(v) with an input value in a set $v \in V = \{0, 1\}$. When a process invokes *propose*(v), we say that the process *proposes* v . The operation returns either a value $v' \in V$ (in which case we say that the process *decides* v') or a special value *abort* $\notin V$ (in which case we say that the invocation *aborts*). A process can invoke the *propose* operation multiple times.

The following properties must be met:

- Validity: every decided value is a proposed value.
- Agreement: no two processes decide differently.
- Obstruction-free termination:
 - If a correct process proposes, it eventually decides or aborts.
 - If a correct process decides, no correct process aborts infinitely often.
 - If there is a time after which *exactly one* correct process p proposes a value sufficiently many times, p eventually decides.

2 Concurrent environment

The goal of the project is to implement OFC for the following environment:

- We have N asynchronous processes. Every process has a distinct *identifier*. The identifiers are publicly known.
- Every two processes can communicate via a reliable asynchronous point-to-point channel.
- Up to $f < N/2$ of the processes are subject to crash failures: a faulty process prematurely stops taking steps of its algorithm. A process that never crashes is called correct.

3 Prerequisites

The project assumes a basic knowledge of Java. Get familiarized with the Java version of AKKA, an actor-based programming model <https://akka.io/docs/>. Check basic constructions in to see how to create an actor, and make the actors communicate.

Check <https://github.com/remisharrock/SLR210Patterns> for sample AKKA patterns which you might want to use.

4 Formalities

The project is pursued in teams of two students.

The implemented system should be provided with a short report describing how the system operates and containing correctness arguments. The team should also prepare a short presentation to be given at the end of the course.

The first project meeting on 26/04 will contain a tutorial on the AKKA programming environment and a discussion of system bootstrapping. The meeting on 17/05 will be used for discussing potential issues and problems. The final meeting on 21/06 will be used for project presentations.

5 Implementation

The implementation should extend the basic construction creating a system of a given size and ensure all-to-all connectivity. Create N actors (processes), and pass references of all N processes to each of them.¹

Use the name `Process` for the process class. For the `Process` class, create methods for invoking the operation *propose*, processing received messages, and returning response indications.

To test the implementation and measure its performance, use the following procedure.

The `main` method selects f processes at random (e.g., using the `shuffle` method from `java.collections`) and sends each of them a special *crash* message. If a process receives a crash message it enters the *fault-prone* mode: for any processed *event* in the algorithm, the process decides, with a fixed probability, if it going to *crash*. If it crashes, it enters the *silent* mode, not reacting to any future event.

For every process, the `main` method then sends a special *launch* message. Once process i receives a launch message, it picks an input value, randomly chosen in $\{0,1\}$ and invokes instances of *propose* operation with this value until a value is decided. (As a basis, one can use the OFC pseudocode to be discussed in the lecture of May 14.)

Use the `LoggingAdapter` class to log both the timing of the invocation and the response of every operation each process performs.

- Emulate a leader election mechanism: after a fixed timeout t_{le} , the `main` method randomly picks up a process that is not *fault-prone* and sends a *hold* message to every other process. After receiving a *hold* message, a process stops invoking *propose* operations.

For example, by invoking `Thread.sleep(50)`, the `main` method “freezes” for 5ms.

An alternative method consists in using the scheduler. For example, the following command:

```
system.scheduler().scheduleOnce(Duration.create(50, TimeUnit.MILLISECONDS),
testActor, "foo", system.dispatcher(), null);
```

results in a message ‘foo’ sent by the scheduler to `testActor` in 50ms.

- Perform the experiment for $N = 3, 10, 100$ (with $f = 1, 4, 49$, respectively) for different values of $t_{le} = 0.5s, 1s, 1.5s, 2s$. For each configuration, measure the time when the first process decides.

Each experiments should be repeated 5 times and the average latency should be evaluated. Build a plot relating the latency with t_{le} for different system sizes.

6 Report

Prepare a short report (up to 15 pages), preferably in English (can also be written in French if English does not feel comfortable). The report should contain:

- The statement of the problem that your implementation solves;
- A high level description of the implementation;

¹Check <https://gitlab.telecom-paristech.fr/petr.kuznetsov/slr210-projects.git> for a example of system creation (`SystemCreation`). Also, the repository contains a program implementing a multi-reader multi-writer atomic register (`AtomicRegister`) in which every process performs a series of *read* and *write* operations on the implemented register.

- A report on performance analysis.

The report and the code of the implementation should be uploaded through the eCampus system by **June 12**.

7 Presentation

The presentation (7 mins) should contain a brief overview of the main features of the algorithm, its correctness arguments and performance. We envision 10 minutes per team (including 3 minutes for questions), so the time bounds are strict.